

```
# Importing all the Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns',None)
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# Lets see what is in the Data
df=pd.read_csv(r"C:\Users\neeth\OneDrive\文档\5th sem\AIML\PROJECT\Clean_Dataset.csv")
df.head()
```

	Unnamed: 0	airline	flight	source_city	departure_time	stops	\
0	0	SpiceJet	SG-8709	Delhi	Evening	zero	
1	1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	
2	2	AirAsia	I5-764	Delhi	Early_Morning	zero	
3	3	Vistara	UK-995	Delhi	Morning	zero	
4	4	Vistara	UK-963	Delhi	Morning	zero	

	arrival_time	destination_city	class	duration	days_left	price
0	Night	Mumbai	Economy	2.17	1	5953
1	Morning	Mumbai	Economy	2.33	1	5953
2	Early_Morning	Mumbai	Economy	2.17	1	5956
3	Afternoon	Mumbai	Economy	2.25	1	5955
4	Morning	Mumbai	Economy	2.33	1	5955

```
# Dropping the useless column 'Unnamed: 0'
df=df.drop('Unnamed: 0',axis=1)
```

```
# A Quick Information about the Data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300153 entries, 0 to 300152
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   airline                300153 non-null object
1   flight                 300153 non-null object
```

```

2  source_city      300153 non-null object
3  departure_time   300153 non-null object
4  stops           300153 non-null object
5  arrival_time     300153 non-null object
6  destination_city 300153 non-null object
7  class            300153 non-null object
8  duration         300153 non-null float64
9  days_left        300153 non-null int64
10 price           300153 non-null int64
dtypes: float64(1), int64(2), object(8)
memory usage: 25.2+ MB

```

Stastical Description of Data
df.describe()

	duration	days_left	price
count	300153.000000	300153.000000	300153.000000
mean	12.221021	26.004751	20889.660523
std	7.191997	13.561004	22697.767366
min	0.830000	1.000000	1105.000000
25%	6.830000	15.000000	4783.000000
50%	11.250000	26.000000	7425.000000
75%	16.170000	38.000000	42521.000000
max	49.830000	49.000000	123071.000000

Size of the data
df.shape

```
(300153, 11)
```

```

df1=df.groupby(['flight','airline'],as_index=False).count()
df1.airline.value_counts()

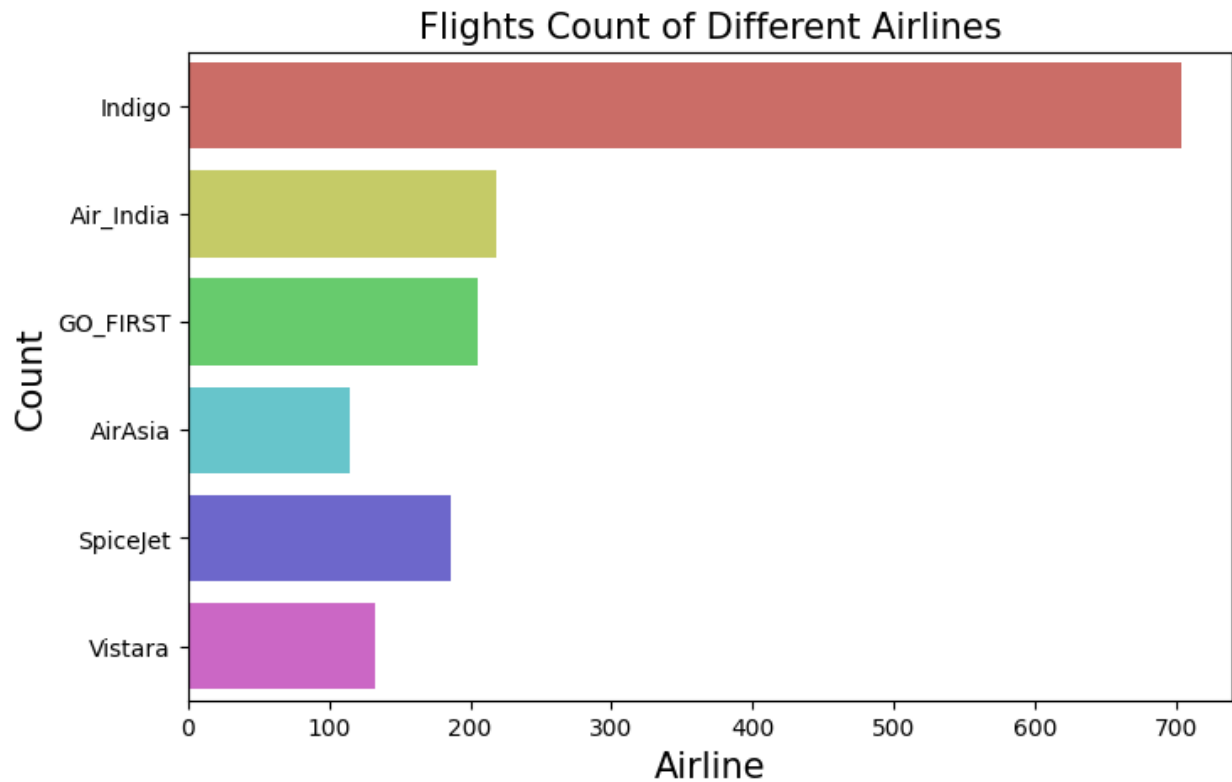
```

```

airline
Indigo      704
Air_India   218
GO_FIRST    205
SpiceJet    186
Vistara     133
AirAsia     115
Name: count, dtype: int64

```

#Indigo becaming as a most popular Airline
plt.figure(figsize=(8,5))
sns.countplot(df1['airline'],palette='hls')
plt.title('Flights Count of Different Airlines',fontsize=15)
plt.xlabel('Airline',fontsize=15)
plt.ylabel('Count',fontsize=15)
plt.show()



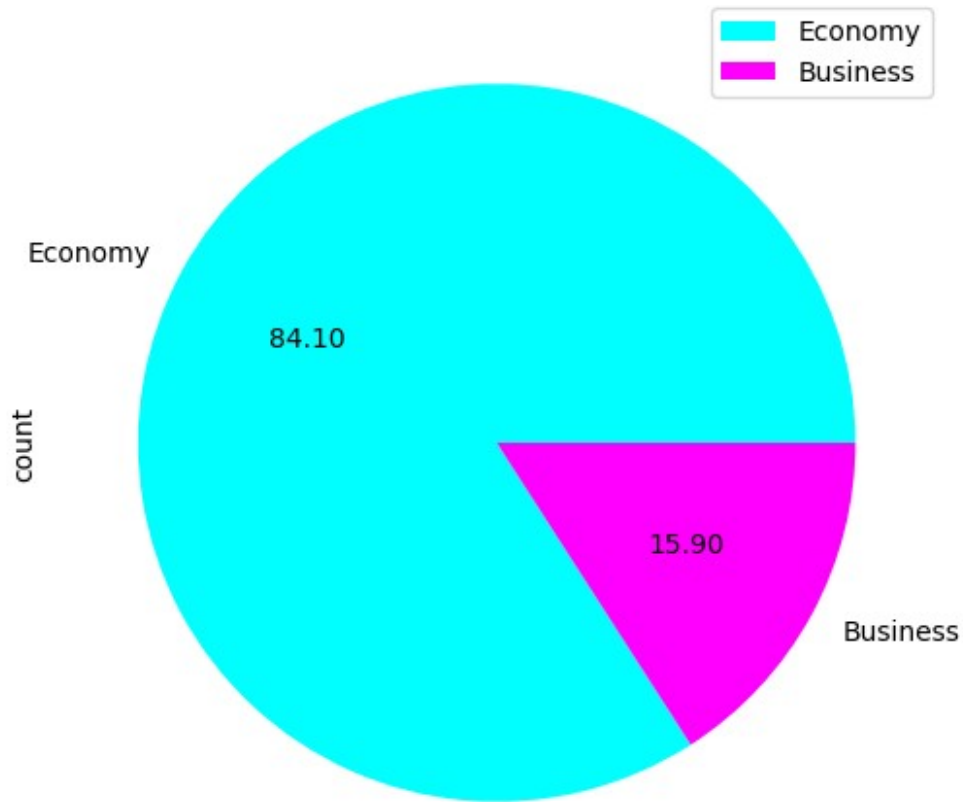
```
df2=df.groupby(['flight','airline','class'],as_index=False).count()  
df2['class'].value_counts()
```

```
class  
Economy    1560  
Business    295  
Name: count, dtype: int64
```


#Most of the Airlines has Economic Class as common

```
plt.figure(figsize=(8,6))  
df2['class'].value_counts().plot(kind='pie',textprops={'color':'black'},  
autopct='%.2f',cmap='cool')  
plt.title('Classes of Different Airlines',fontsize=15)  
plt.legend(['Economy','Business'])  
plt.show()
```


Classes of Different Airlines



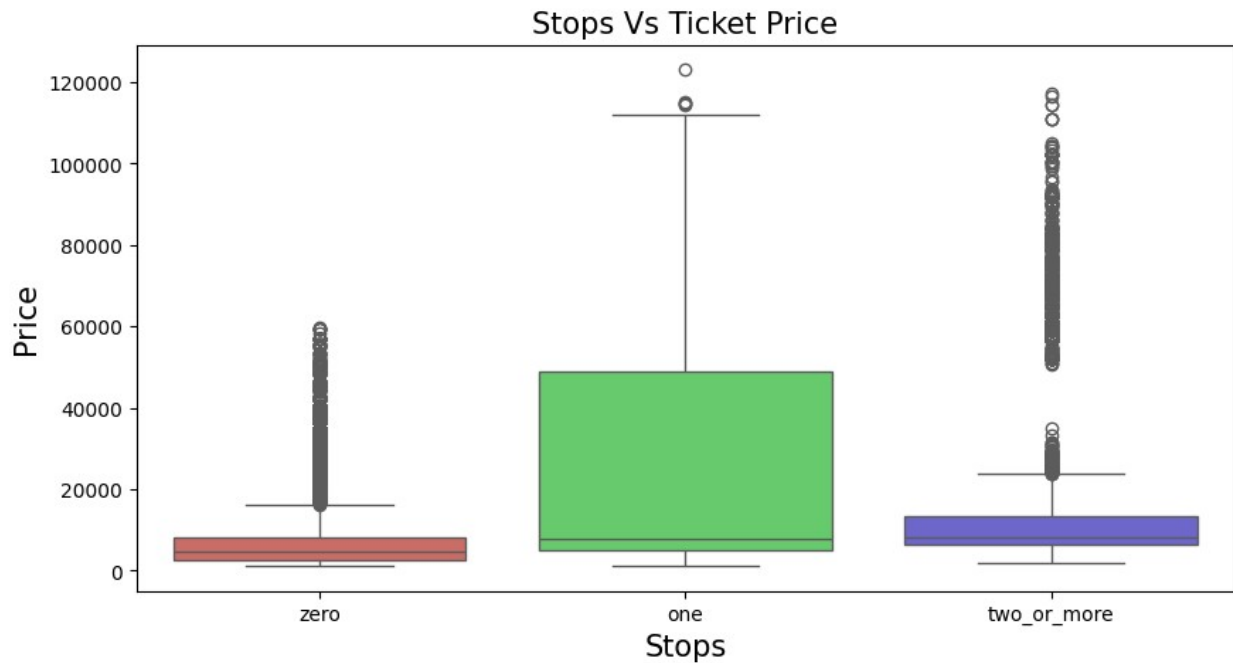
```
plt.figure(figsize=(15,5))
sns.boxplot(x=df['airline'],y=df['price'],palette='hls')
plt.title('Airlines Vs Price',fontsize=15)
plt.xlabel('Airline',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```




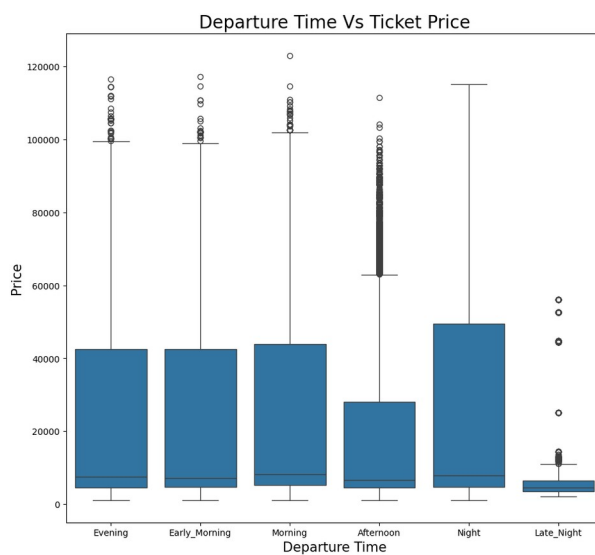
```
plt.figure(figsize=(10,5))
sns.boxplot(x='class',y='price',data=df,palette='hls')
plt.title('Class Vs Ticket Price',fontsize=15)
plt.xlabel('Class',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```



```
plt.figure(figsize=(10,5))
sns.boxplot(x='stops',y='price',data=df,palette='hls')
plt.title('Stops Vs Ticket Price',fontsize=15)
plt.xlabel('Stops',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```




```
plt.figure(figsize=(24,10))
plt.subplot(1,2,1)
sns.boxplot(x='departure_time',y='price',data=df)
plt.title('Departure Time Vs Ticket Price',fontsize=20)
plt.xlabel('Departure Time',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.subplot(1,2,2)
sns.boxplot(x='arrival_time',y='price',data=df,palette='hls')
plt.title('Arrival Time Vs Ticket Price',fontsize=20)
plt.xlabel('Arrival Time',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```



```

plt.figure(figsize=(24,10))
plt.subplot(1,2,1)
sns.boxplot(x='source_city',y='price',data=df)
plt.title('Source City Vs Ticket Price',fontsize=20)
plt.xlabel('Source City',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.subplot(1,2,2)
sns.boxplot(x='destination_city',y='price',data=df,palette='hls')
plt.title('Destination City Vs Ticket Price',fontsize=20)
plt.xlabel('Destination City',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()

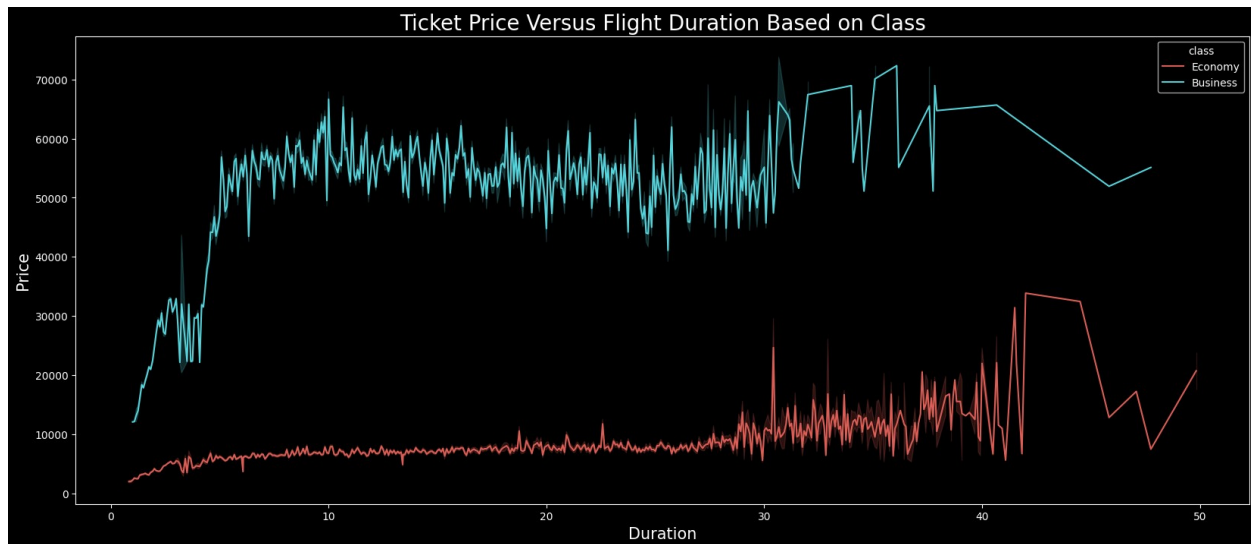
```




```

plt.style.use('dark_background')
plt.figure(figsize=(20,8))
sns.lineplot(data=df,x='duration',y='price',hue='class',palette='hls')
plt.title('Ticket Price Versus Flight Duration Based on
Class',fontsize=20)
plt.xlabel('Duration',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()

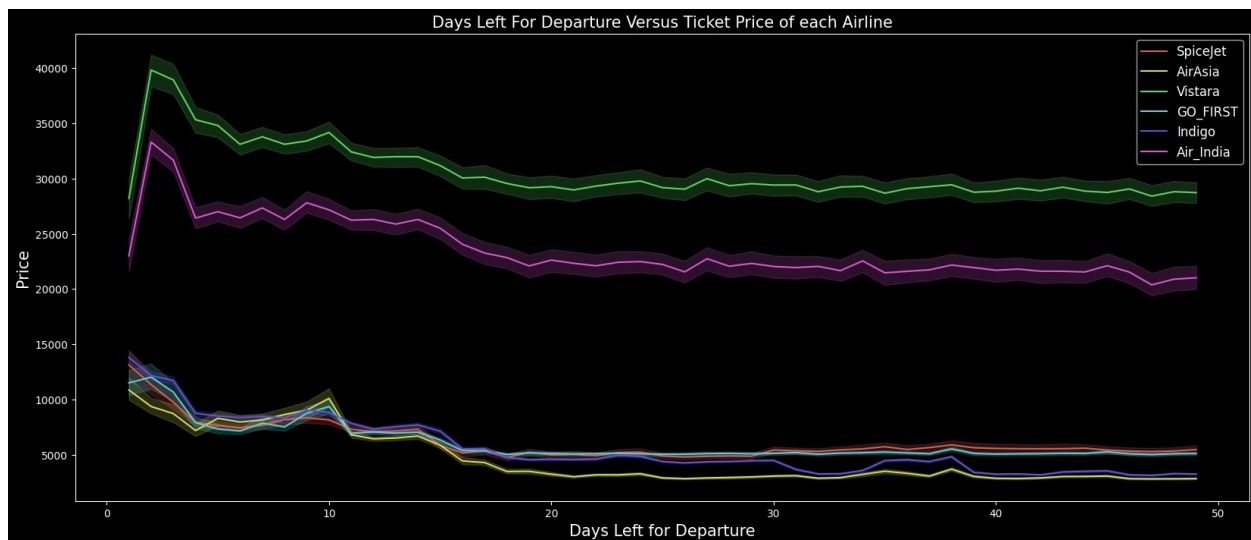
```



```
plt.figure(figsize=(20,8))
sns.lineplot(data=df,x='days_left',y='price',color='blue')
plt.title('Days Left For Departure Versus Ticket Price',fontsize=20)
plt.xlabel('Days Left for Departure',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```



```
plt.figure(figsize=(20,8))
sns.lineplot(data=df,x='days_left',y='price',color='blue',hue='airline',
            palette='hls')
plt.title('Days Left For Departure Versus Ticket Price of each Airline',
        fontsize=15)
plt.legend(fontsize=12)
plt.xlabel('Days Left for Departure',fontsize=15)
plt.ylabel('Price',fontsize=15)
plt.show()
```

```
df.groupby(['flight', 'source_city', 'destination_city', 'airline', 'class'], as_index=False).count().groupby(['source_city', 'destination_city'], as_index=False)['flight'].count().head(10)
```

	source_city	destination_city	flight
0	Bangalore	Chennai	106
1	Bangalore	Delhi	227
2	Bangalore	Hyderabad	132
3	Bangalore	Kolkata	171
4	Bangalore	Mumbai	175
5	Chennai	Bangalore	69
6	Chennai	Delhi	105
7	Chennai	Hyderabad	82
8	Chennai	Kolkata	110
9	Chennai	Mumbai	94

```
df.groupby(['airline', 'source_city', 'destination_city'], as_index=False)['price'].mean().head(10)
```

	airline	source_city	destination_city	price
0	AirAsia	Bangalore	Chennai	2073.043478
1	AirAsia	Bangalore	Delhi	4807.092426
2	AirAsia	Bangalore	Hyderabad	2931.494792
3	AirAsia	Bangalore	Kolkata	4443.468160
4	AirAsia	Bangalore	Mumbai	3342.385350
5	AirAsia	Chennai	Bangalore	1914.760870
6	AirAsia	Chennai	Delhi	3697.314003
7	AirAsia	Chennai	Hyderabad	2053.182540
8	AirAsia	Chennai	Kolkata	3682.338762
9	AirAsia	Chennai	Mumbai	2691.100000

```
# Creating a Back up File
df_bk=df.copy()
```

```

# Converting the labels into a numeric form using Label Encoder
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for col in df.columns:
    if df[col].dtype=='object':
        df[col]=le.fit_transform(df[col])

# storing the Dependent Variables in X and Independent Variable in Y
x=df.drop(['price'],axis=1)
y=df['price']

# Splitting the Data into Training set and Testing Set
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=42)
x_train.shape,x_test.shape,y_train.shape,y_test.shape

((210107, 10), (90046, 10), (210107,), (90046,))

# Scaling the values to convert the int values to Machine Languages
from sklearn.preprocessing import MinMaxScaler
mmScaler=MinMaxScaler(feature_range=(0,1))
x_train=mmScaler.fit_transform(x_train)
x_test=mmScaler.fit_transform(x_test)
x_train=pd.DataFrame(x_train)
x_test=pd.DataFrame(x_test)

a={'Model Name':[], 'Mean_Absolute_Error_MAE':[] , 'Adj_R_Square':
[] , 'Root_Mean_Squared_Error_RMSE':
[] , 'Mean_Absolute_Percentage_Error_MAPE':
[] , 'Mean_Squared_Error_MSE':[] , 'Root_Mean_Squared_Log_Error_RMSLE':
[] , 'R2_score':[]}
Results=pd.DataFrame(a)
Results.head()

Empty DataFrame
Columns: [Model Name, Mean_Absolute_Error_MAE, Adj_R_Square,
Root_Mean_Squared_Error_RMSE, Mean_Absolute_Percentage_Error_MAPE,
Mean_Squared_Error_MSE, Root_Mean_Squared_Log_Error_RMSLE, R2_score]
Index: []

# Build the Regression / Regressor models

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn import linear_model
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
import xgboost as xgb
from sklearn.neighbors import KNeighborsRegressor

```

```

from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import BaggingRegressor
from sklearn.ensemble import GradientBoostingRegressor

# Create objects of Regression / Regressor models with default hyper-
parameters

modelmlg = LinearRegression()
modeldcr = DecisionTreeRegressor()
modelbag = BaggingRegressor()
modelrfr = RandomForestRegressor()
modelSVR = SVR()
modelXGR = xgb.XGBRegressor()
modelKNN = KNeighborsRegressor(n_neighbors=5)
modelETR = ExtraTreesRegressor()
modelRE=Ridge()
modelL0=linear_model.Lasso(alpha=0.1)
modelGBR = GradientBoostingRegressor(loss='squared_error',
learning_rate=0.1, n_estimators=100, subsample=1.0,
criterion='friedman_mse',
min_samples_split=2, min_samples_leaf=1,
min_weight_fraction_leaf=0.0,
max_depth=3, min_impurity_decrease=0.0,
init=None, random_state=None,
max_features=None,
alpha=0.9, verbose=0,
max_leaf_nodes=None, warm_start=False,
validation_fraction=0.1,
n_iter_no_change=None, tol=0.0001, ccp_alpha=0.0)

# Evaluation matrix for all the algorithms

MM = [modelmlg, modeldcr, modelrfr, modelKNN, modelETR, modelGBR,
modelXGR, modelbag,modelRE,modelL0]

for models in MM:

    # Fit the model with train data

    models.fit(x_train, y_train)

    # Predict the model with test data

    y_pred = models.predict(x_test)

    # Print the model name

    print('Model Name: ', models)

    # Evaluation metrics for Regression analysis

```

```

from sklearn import metrics

print('Mean Absolute Error (MAE):',
round(metrics.mean_absolute_error(y_test, y_pred),3))
print('Mean Squared Error (MSE):',
round(metrics.mean_squared_error(y_test, y_pred),3))
print('Root Mean Squared Error (RMSE):',
round(np.sqrt(metrics.mean_squared_error(y_test, y_pred)),3))
print('R2_score:', round(metrics.r2_score(y_test, y_pred),6))
print('Root Mean Squared Log Error (RMSLE):',
round(np.log(np.sqrt(metrics.mean_squared_error(y_test, y_pred))),3))

# Define the function to calculate the MAPE - Mean Absolute
Percentage Error

def MAPE (y_test, y_pred):
    y_test, y_pred = np.array(y_test), np.array(y_pred)
    return np.mean(np.abs((y_test - y_pred) / y_test)) * 100

# Evaluation of MAPE

result = MAPE(y_test, y_pred)
print('Mean Absolute Percentage Error (MAPE):', round(result, 2),
'%%')

# Calculate Adjusted R squared values

r_squared = round(metrics.r2_score(y_test, y_pred),6)
adjusted_r_squared = round(1 - (1-r_squared)*(len(y)-1)/(len(y) -
x.shape[1]-1),6)
print('Adj R Square: ', adjusted_r_squared)

print('-----')
-----

#-----
-----

new_row = {'Model Name' : models,
           'Mean_Absolute_Error_MAE' :
metrics.mean_absolute_error(y_test, y_pred),
           'Adj_R_Square' : adjusted_r_squared,
           'Root_Mean_Squared_Error_RMSE' :
np.sqrt(metrics.mean_squared_error(y_test, y_pred)),
           'Mean_Absolute_Percentage_Error_MAPE' : result,
           'Mean_Squared_Error_MSE' :
metrics.mean_squared_error(y_test, y_pred),
           'Root_Mean_Squared_Log_Error_RMSLE':
np.log(np.sqrt(metrics.mean_squared_error(y_test, y_pred))),
           'R2_score' : metrics.r2_score(y_test, y_pred)}

```

```
import pandas as pd
Results = pd.concat([Results, pd.DataFrame([new_row])],
ignore_index=True)
```

```
#-----
```

```
Model Name: LinearRegression()
Mean Absolute Error (MAE): 4630.296
Mean Squared Error (MSE): 49070241.265
Root Mean Squared Error (RMSE): 7005.015
R2_score: 0.904656
Root Mean Squared Log Error (RMSLE): 8.854
Mean Absolute Percentage Error (MAPE): 43.89 %
Adj R Square: 0.904653
```

```
-----
Model Name: DecisionTreeRegressor()
Mean Absolute Error (MAE): 1270.674
Mean Squared Error (MSE): 14012980.306
Root Mean Squared Error (RMSE): 3743.392
R2_score: 0.972773
Root Mean Squared Log Error (RMSLE): 8.228
Mean Absolute Percentage Error (MAPE): 8.66 %
Adj R Square: 0.972772
```

```
-----
Model Name: RandomForestRegressor()
Mean Absolute Error (MAE): 1172.282
Mean Squared Error (MSE): 8228203.693
Root Mean Squared Error (RMSE): 2868.485
R2_score: 0.984013
Root Mean Squared Log Error (RMSLE): 7.962
Mean Absolute Percentage Error (MAPE): 7.94 %
Adj R Square: 0.984012
```

```
-----
Model Name: KNeighborsRegressor()
Mean Absolute Error (MAE): 1854.896
Mean Squared Error (MSE): 14590102.958
Root Mean Squared Error (RMSE): 3819.699
R2_score: 0.971651
Root Mean Squared Log Error (RMSLE): 8.248
Mean Absolute Percentage Error (MAPE): 11.13 %
Adj R Square: 0.97165
```

```
-----
Model Name: ExtraTreesRegressor()
Mean Absolute Error (MAE): 1139.641
Mean Squared Error (MSE): 7773122.609
Root Mean Squared Error (RMSE): 2788.032
```

R2_score: 0.984897
Root Mean Squared Log Error (RMSLE): 7.933
Mean Absolute Percentage Error (MAPE): 7.59 %
Adj R Square: 0.984896

Model Name: GradientBoostingRegressor()
Mean Absolute Error (MAE): 2808.069
Mean Squared Error (MSE): 22221132.303
Root Mean Squared Error (RMSE): 4713.93
R2_score: 0.956824
Root Mean Squared Log Error (RMSLE): 8.458
Mean Absolute Percentage Error (MAPE): 20.67 %
Adj R Square: 0.956823

Model Name: XGBRegressor(base_score=None, booster=None,
callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, device=None,
early_stopping_rounds=None,
enable_categorical=False, eval_metric=None,
feature_types=None,
gamma=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=None,
max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan,
monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)
Mean Absolute Error (MAE): 1862.407
Mean Squared Error (MSE): 11680762.0
Root Mean Squared Error (RMSE): 3417.713
R2_score: 0.977304
Root Mean Squared Log Error (RMSLE): 8.137
Mean Absolute Percentage Error (MAPE): 14.23 %
Adj R Square: 0.977303

Model Name: BaggingRegressor()
Mean Absolute Error (MAE): 1207.99
Mean Squared Error (MSE): 8966813.523
Root Mean Squared Error (RMSE): 2994.464
R2_score: 0.982577
Root Mean Squared Log Error (RMSLE): 8.005
Mean Absolute Percentage Error (MAPE): 8.25 %
Adj R Square: 0.982576

```

-----
Model Name: Ridge()
Mean Absolute Error (MAE): 4630.313
Mean Squared Error (MSE): 49070187.277
Root Mean Squared Error (RMSE): 7005.012
R2_score: 0.904656
Root Mean Squared Log Error (RMSLE): 8.854
Mean Absolute Percentage Error (MAPE): 43.89 %
Adj R Square: 0.904653
-----

```

```

-----
Model Name: Lasso(alpha=0.1)
Mean Absolute Error (MAE): 4630.179
Mean Squared Error (MSE): 49070111.63
Root Mean Squared Error (RMSE): 7005.006
R2_score: 0.904656
Root Mean Squared Log Error (RMSLE): 8.854
Mean Absolute Percentage Error (MAPE): 43.89 %
Adj R Square: 0.904653
-----

```

Results

	Model Name
Mean_Absolute_Error_MAE \	
0	LinearRegression()
4630.295614	
1	DecisionTreeRegressor()
1267.046858	
2 (DecisionTreeRegressor(max_features=1.0, rando...	
1171.195628	
3	KNeighborsRegressor()
1854.469527	
4 (ExtraTreeRegressor(random_state=370408196), E...	
1148.979440	
5 ([DecisionTreeRegressor(criterion='friedman_ms...	
2808.074076	
6 XGBRegressor(base_score=None, booster=None, ca...	
1862.407227	
7 (DecisionTreeRegressor(random_state=510902622)...	
1210.905216	
8	Ridge()
4630.313301	
9	Lasso(alpha=0.1)
4630.179207	
Adj_R_Square	Root_Mean_Squared_Error_RMSE \
0	0.904653 7005.015436

1	0.972946	3731.414789
2	0.984113	2859.347334
3	0.971645	3820.074173
4	0.984726	2803.632993
5	0.956823	4713.930286
6	0.977303	3417.712978
7	0.982816	2973.788019
8	0.904653	7005.011583
9	0.904653	7005.006183

	Mean_Absolute_Percentage_Error_MAPE	Mean_Squared_Error_MSE	\
0	43.888567	4.907024e+07	
1	8.636776	1.392346e+07	
2	7.910083	8.175867e+06	
3	11.113824	1.459297e+07	
4	7.629837	7.860358e+06	
5	20.673080	2.222114e+07	
6	14.232981	1.168076e+07	
7	8.290388	8.843415e+06	
8	43.888754	4.907019e+07	
9	43.885459	4.907011e+07	

	Root_Mean_Squared_Log_Error_RMSLE	R2_score
0	8.854382	0.904656
1	8.224543	0.972947
2	7.958349	0.984114
3	8.248025	0.971646
4	7.938671	0.984727
5	8.458277	0.956824
6	8.136727	0.977304
7	7.997592	0.982817
8	8.854381	0.904656
9	8.854380	0.904656

```
models=['LinearRegression','DecisionTreeRegressor','RandomForestRegressor',
'KNeighborsRegressor','ExtraTreesRegressor','GradientBoostingRegressor',
'XGBRegressor','BaggingRegressor','Ridge Regression','Lasso Regression']
result=pd.DataFrame({'Model_Name':models})
result['Adj_R_Square']=Results['Adj_R_Square']
result['Mean_Absolute_Error_MAE']=Results['Mean_Absolute_Error_MAE']
result['Root_Mean_Squared_Error_RMSE']=Results['Root_Mean_Squared_Error_RMSE']
result['Mean_Absolute_Percentage_Error_MAPE']=Results['Mean_Absolute_Percentage_Error_MAPE']
result['Mean_Squared_Error_MSE']=Results['Mean_Squared_Error_MSE']
result['Root_Mean_Squared_Log_Error_RMSLE']=Results['Root_Mean_Squared_Log_Error_RMSLE']
result['R2_score']=Results['R2_score']
result=result.sort_values(by='Adj_R_Square',ascending=False).reset_index
```



```
ex(drop=True)
result
```

	Model_Name	Adj_R_Square	Mean_Absolute_Error_MAE	\
0	ExtraTreesRegressor	0.984726	1148.979440	
1	RandomForestRegressor	0.984113	1171.195628	
2	BaggingRegressor	0.982816	1210.905216	
3	XGBRegressor	0.977303	1862.407227	
4	DecisionTreeRegressor	0.972946	1267.046858	
5	KNeighborsRegressor	0.971645	1854.469527	
6	GradientBoostingRegressor	0.956823	2808.074076	
7	LinearRegression	0.904653	4630.295614	
8	Ridge Regression	0.904653	4630.313301	
9	Lasso Regression	0.904653	4630.179207	

	Root_Mean_Squared_Error_RMSE	Mean_Absolute_Percentage_Error_MAPE	\
0	2803.632993	7.629837	
1	2859.347334	7.910083	
2	2973.788019	8.290388	
3	3417.712978	14.232981	
4	3731.414789	8.636776	
5	3820.074173	11.113824	
6	4713.930286	20.673080	
7	7005.015436	43.888567	
8	7005.011583	43.888754	
9	7005.006183	43.885459	

	Mean_Squared_Error_MSE	Root_Mean_Squared_Log_Error_RMSLE	R2_score
0	7.860358e+06	7.938671	0.984727
1	8.175867e+06	7.958349	0.984114
2	8.843415e+06	7.997592	0.982817
3	1.168076e+07	8.136727	0.977304
4	1.392346e+07	8.224543	0.972947
5	1.459297e+07	8.248025	0.971646

6	2.222114e+07	8.458277	0.956824
7	4.907024e+07	8.854382	0.904656
8	4.907019e+07	8.854381	0.904656
9	4.907011e+07	8.854380	0.904656

```
#Trainig the model with
modelETR.fit(x_train, y_train)
```

```
# Predict the model with test data
```

```
y_pred = modelETR.predict(x_test)
```

```
out=pd.DataFrame({'Price_actual':y_test,'Price_pred':y_pred})
result=df_bk.merge(out,left_index=True,right_index=True)
```

```
result.sample(10)
```

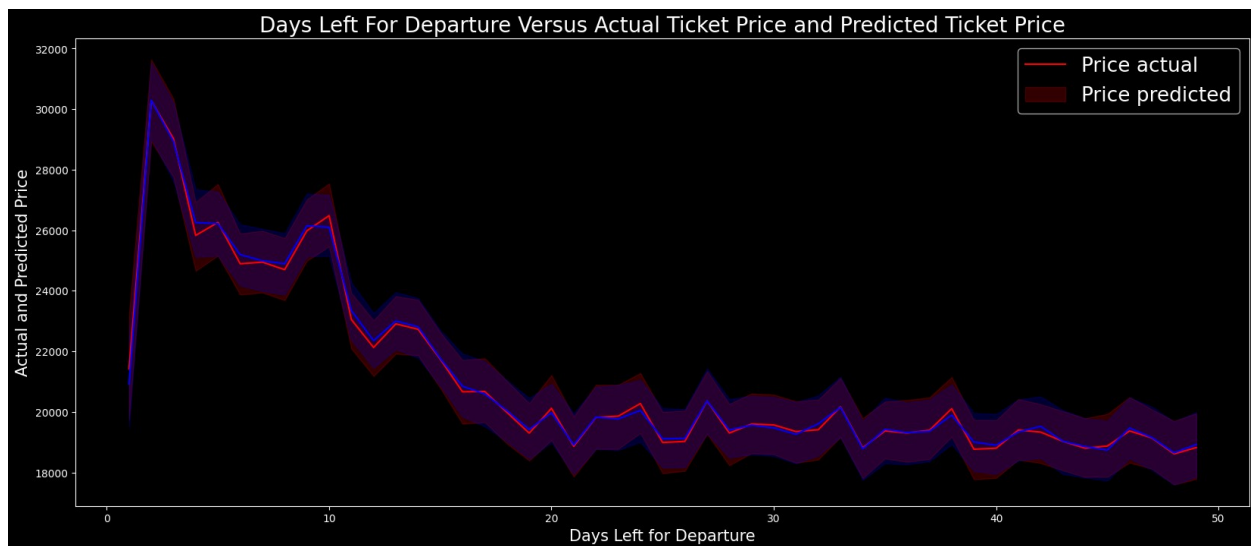
	airline	flight	source_city	departure_time	stops
arrival_time \					
172088	GO_FIRST	G8-123	Hyderabad	Early_Morning	zero
Early_Morning					
172328	AirAsia	I5-510	Hyderabad	Early_Morning	one
Afternoon					
267403	Vistara	UK-772	Kolkata	Morning	one
Morning					
151587	Indigo	6E-345	Kolkata	Morning	one
Evening					
226901	Vistara	UK-841	Mumbai	Morning	one
Evening					
157140	Indigo	6E-7156	Hyderabad	Afternoon	one
Night					
263123	Air_India	AI-773	Kolkata	Evening	one
Morning					
260667	Vistara	UK-864	Bangalore	Evening	one
Evening					
26986	AirAsia	I5-721	Delhi	Night	one
Morning					
157987	Vistara	UK-874	Hyderabad	Morning	one
Late_Night					

	destination_city	class	duration	days_left	price
Price_actual \					
172088	Kolkata	Economy	1.83	15	5292
5292					
172328	Kolkata	Economy	6.00	17	2056
2056					

267403	Bangalore	Business	22.42	3	75023
75023					
151587	Chennai	Economy	10.00	39	3015
3015					
226901	Delhi	Business	5.67	20	60813
60813					
157140	Delhi	Economy	6.17	39	3243
3243					
263123	Delhi	Business	17.50	34	53743
53743					
260667	Chennai	Business	24.92	38	44280
44280					
26986	Kolkata	Economy	12.25	40	3175
3175					
157987	Delhi	Economy	16.42	45	5656
5656					

	Price_pred
172088	5431.45
172328	3579.41
267403	68566.20
151587	3315.70
226901	54000.05
157140	3669.54
263123	53743.00
260667	44354.27
26986	2621.60
157987	5652.85

```
plt.figure(figsize=(20,8))
sns.lineplot(data=result,x='days_left',y='Price_actual',color='red')
sns.lineplot(data=result,x='days_left',y='Price_pred',color='blue')
plt.title('Days Left For Departure Versus Actual Ticket Price and Predicted Ticket Price',fontsize=20)
plt.legend(labels=['Price actual','Price predicted'],fontsize=19)
plt.xlabel('Days Left for Departure',fontsize=15)
plt.ylabel('Actual and Predicted Price',fontsize=15)
plt.show()
```



```
plt.figure(figsize=(10,5))
sns.regplot(x='Price_actual',y='Price_pred',data=result,color='cyan')
plt.title('Actual Price Vs Predicted Price ',fontsize=20)
plt.xlabel('Actual Price',fontsize=15)
plt.ylabel('Predicted Price',fontsize=15)
plt.show()
```



```
import pandas as pd
import numpy as np
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.linear_model import LinearRegression # Import
```

LinearRegression

```
# Initialize Results DataFrame
Results = pd.DataFrame(columns=[
    'Model Name', 'Mean_Absolute_Error_MAE', 'Adj_R_Square',
    'Root_Mean_Squared_Error_RMSE',
    'Mean_Absolute_Percentage_Error_MAPE',
    'Mean_Squared_Error_MSE', 'Root_Mean_Squared_Log_Error_RMSLE',
    'R2_score'])

# Iterate through each model in the MM list
for models in MM:
    # Check if the current model is LinearRegression
    if isinstance(models, LinearRegression):
        # Fit the model with train data
        models.fit(x_train, y_train)

        # Predict the model with test data
        y_pred = models.predict(x_test)

        # Model evaluation
        mae = mean_absolute_error(y_test, y_pred)
        mse = mean_squared_error(y_test, y_pred)
        rmse = np.sqrt(mse)
        r2 = r2_score(y_test, y_pred)

        # Calculate MAPE (Mean Absolute Percentage Error)
        def MAPE(y_test, y_pred):
            y_test, y_pred = np.array(y_test), np.array(y_pred)
            return np.mean(np.abs((y_test - y_pred) / y_test)) * 100

        mape = MAPE(y_test, y_pred)

        # Adjusted R-Square calculation
        adj_r2 = 1 - (1 - r2) * (len(y_test) - 1) / (len(y_test) -
x_test.shape[1] - 1)

        # RMSLE calculation
        rmsle = np.log(np.sqrt(mse))

        # Print evaluation metrics
        print(f"Model Name: {models}")
        print(f"Mean Absolute Error (MAE): {mae:.3f}")
        print(f"Mean Squared Error (MSE): {mse:.3f}")
        print(f"Root Mean Squared Error (RMSE): {rmse:.3f}")
        print(f"R2 Score: {r2:.6f}")
        print(f"Mean Absolute Percentage Error (MAPE): {mape:.2f}%")
        print(f"Root Mean Squared Log Error (RMSLE): {rmsle:.3f}")
        print("-" * 100)
```

```

# Prepare new row as a DataFrame
new_row = pd.DataFrame({
    'Model Name': [str(models)],
    'Mean_Absolute_Error_MAE': [mae],
    'Adj_R_Square': [adj_r2],
    'Root_Mean_Squared_Error_RMSE': [rmse],
    'Mean_Absolute_Percentage_Error_MAPE': [mape],
    'Mean_Squared_Error_MSE': [mse],
    'Root_Mean_Squared_Log_Error_RMSLE': [rmsle],
    'R2_score': [r2]
})

# Concatenate the new row with the existing Results DataFrame
Results = pd.concat([Results, new_row], ignore_index=True)

```

```

# After the loop, you can check the Results DataFrame
print(Results)

```

```

Model Name: LinearRegression()
Mean Absolute Error (MAE): 4630.296
Mean Squared Error (MSE): 49070241.265
Root Mean Squared Error (RMSE): 7005.015
R2 Score: 0.904656
Mean Absolute Percentage Error (MAPE): 43.89%
Root Mean Squared Log Error (RMSLE): 8.854

```

```

-----
-----
      Model Name  Mean_Absolute_Error_MAE  Adj_R_Square  \
0  LinearRegression()          4630.295614          0.904646

      Root_Mean_Squared_Error_RMSE
Mean_Absolute_Percentage_Error_MAPE  \
0              7005.015436              43.888567

      Mean_Squared_Error_MSE  Root_Mean_Squared_Log_Error_RMSLE  R2_score
0              4.907024e+07              8.854382      0.904656

```

```

from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.naive_bayes import GaussianNB

# Load a dataset (example: Iris dataset)
data = load_iris()
X = data.data # Feature data
y = data.target # Target labels

# Split the dataset into training and testing sets

```

```

x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# List of models to train (in this case, Naive Bayes)
MM = [GaussianNB()]

# Iterate through each model in the MM list
for model in MM:
    # Fit the model with train data
    model.fit(x_train, y_train)

    # Predict the model with test data
    y_pred = model.predict(x_test)

    # Print or evaluate the results (you can add metrics like
    accuracy, etc.)
    print(f'Model: {model.__class__.__name__}')
    print(f'Predictions: {y_pred}')

Model: GaussianNB
Predictions: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 2 2 1 1 2 0 2 0 2 2 2 2
0 0 0 0 1 0 0 2 1
0 0 0 2 1 1 0 0]

from sklearn.metrics import accuracy_score, confusion_matrix

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print(f'Confusion Matrix:\n{conf_matrix}')

Accuracy: 97.78%
Confusion Matrix:
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split

# Example DataFrame 'result' (replace this with your actual data)
import pandas as pd

result = pd.DataFrame({
    'days_left': [30, 60, 90, 120, 150], # Example days_left

```

```

    'Price_actual': [100, 150, 120, 90, 110] # Example actual prices
})

# Prepare the data
X = result[['days_left']] # Feature: 'days_left'
y = result['Price_actual'] # Target: 'Price_actual'

# Split the dataset into train and test sets
x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Fit the Linear Regression model
lr_model = LinearRegression()
lr_model.fit(x_train, y_train)
y_pred_lr = lr_model.predict(x_test)

# Fit the Naive Bayes model (GaussianNB is generally used for
classification)
nb_model = GaussianNB()
y_train_nb = y_train.astype(int) # Naive Bayes expects integer labels
nb_model.fit(x_train, y_train_nb)
y_pred_nb = nb_model.predict(x_test)

# Plotting the results
plt.figure(figsize=(20, 8))

# Plot Linear Regression predictions
sns.lineplot(x=x_test['days_left'], y=y_pred_lr, color='blue',
label='Price predicted (Linear Regression)')

# Plot Naive Bayes predictions
sns.lineplot(x=x_test['days_left'], y=y_pred_nb, color='green',
label='Price predicted (Naive Bayes)')

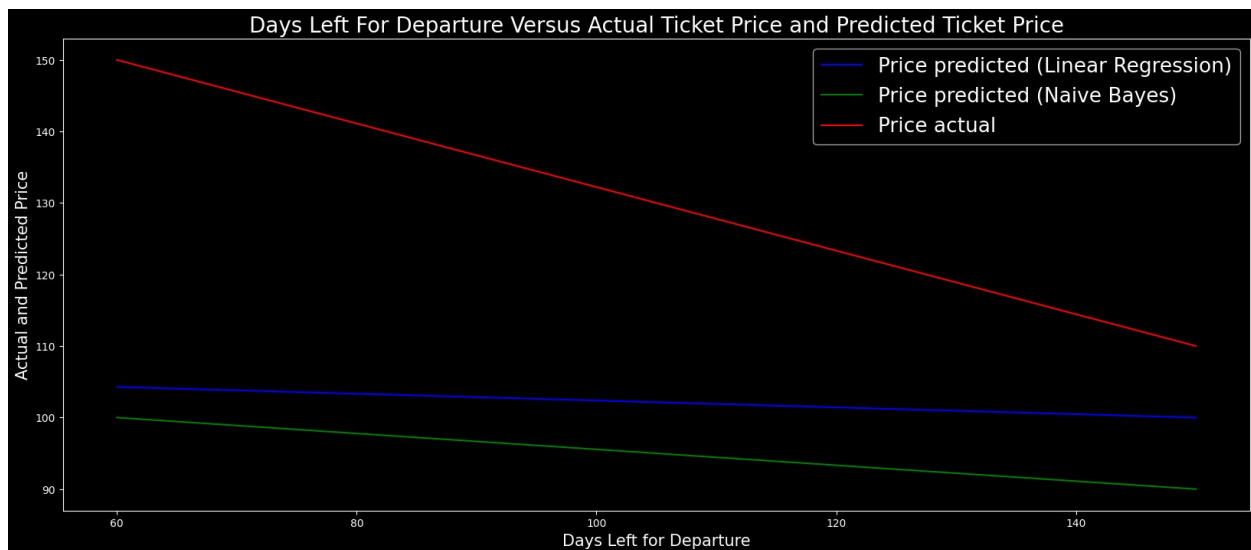
# Plot actual prices for the test set
sns.lineplot(x=x_test['days_left'], y=y_test, color='red',
label='Price actual')

# Title and labels
plt.title('Days Left For Departure Versus Actual Ticket Price and
Predicted Ticket Price', fontsize=20)
plt.xlabel('Days Left for Departure', fontsize=15)
plt.ylabel('Actual and Predicted Price', fontsize=15)

# Show legend
plt.legend(fontsize=19)

# Display the plot
plt.show()

```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split

# Example DataFrame 'result' (replace this with your actual data)
import pandas as pd

result = pd.DataFrame({
    'days_left': [30, 60, 90, 120, 150], # Example days_left
    'Price_actual': [100, 150, 120, 90, 110] # Example actual prices
})

# Prepare the data
X = result[['days_left']] # Feature: 'days_left'
y = result['Price_actual'] # Target: 'Price_actual'

# Split the dataset into train and test sets
x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Fit the Linear Regression model
lr_model = LinearRegression()
lr_model.fit(x_train, y_train)
y_pred_lr = lr_model.predict(x_test)

# Fit the Naive Bayes model (GaussianNB is generally used for
classification)
nb_model = GaussianNB()
y_train_nb = y_train.astype(int) # Naive Bayes expects integer labels
nb_model.fit(x_train, y_train_nb)
y_pred_nb = nb_model.predict(x_test)
```

```
# Add the predictions to the 'result' DataFrame (for comparison)
# Assuming you're interested in the Linear Regression predictions
result['Price_pred'] = lr_model.predict(X)

# Plotting the Actual vs Predicted prices
plt.figure(figsize=(10, 5))
sns.regplot(x='Price_actual', y='Price_pred', data=result,
color='cyan')
plt.title('Actual Price Vs Predicted Price', fontsize=20)
plt.xlabel('Actual Price', fontsize=15)
plt.ylabel('Predicted Price', fontsize=15)
plt.show()
```

